

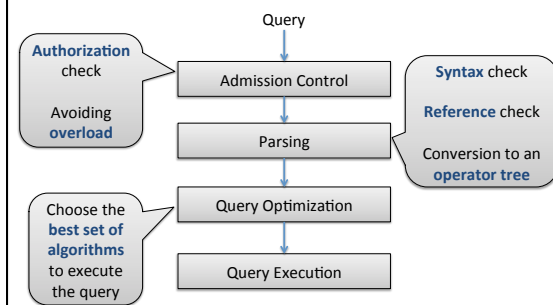
## Introduction to Data Management CSE 344

Extra Lecture: Basics of Query Evaluation  
Presenter: Prasang Upadhyaya

## Outline

- Life of a query in a DBMS
  - Query plans
- Indexing

## Life of a Query



## Query Plan

Output of parsing

A **tree of relational operators**

Joins, Aggregates, Selections, Projections

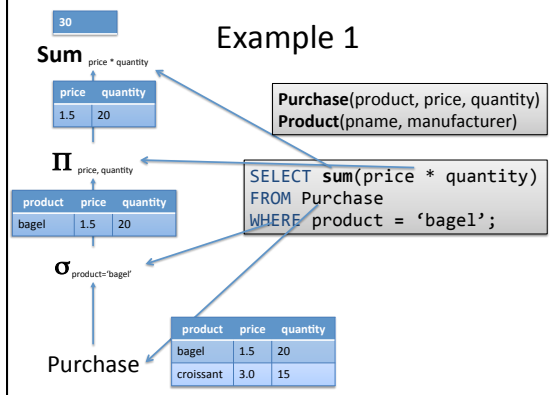
Operators consume relations and produce relations

**Leaves are relations**

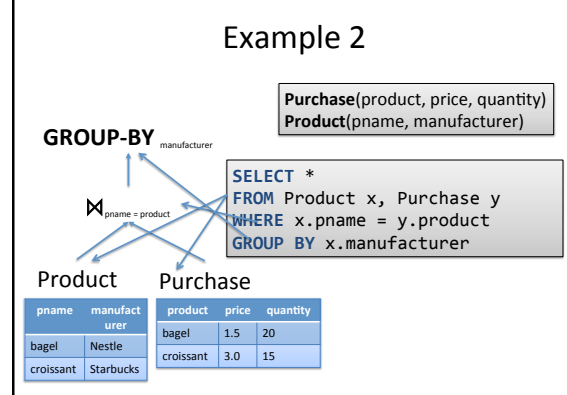
**Output at the root**

Every SQL statement is converted to a Query Plan

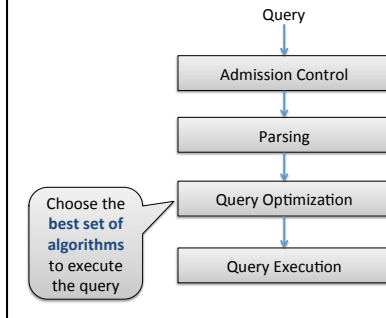
## Example 1



## Example 2



## Life of a Query



## Query Optimization

Find the **best way** to execute a query

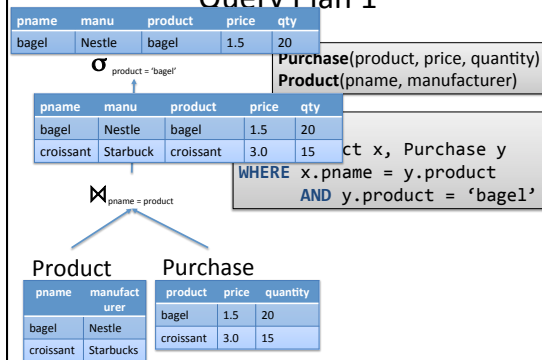
Given a query plan

Construct **alternate, equivalent** plans

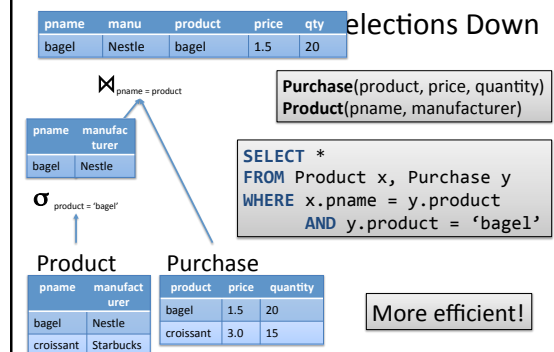
Determine their **cost**

Select the **cheapest** plan

## Query Plan 1



## Selections Down



## Cost of Query Plan

- Used to decide **which plan to pick**
- Cost approximates **how much time** the query would take
- Depends on:
  - How relations are **stored on disk**
  - How the **operator is implemented**

## Data Storage

- DBMSs store data in **files**
- Most common organization is row-wise storage
- On disk, a file is split into blocks
- Each block contains a set of tuples

10	...
20	
30	
40	
50	
60	
70	
80	

In the example, we have 4 blocks with 2 tuples each

## Cost Parameters

- In database systems the data is on disk
- Cost = total number of I/Os**
  - This is a simplification
  - Normally, need to consider IO, CPU, and network
- Cost of an operation = number of disk I/Os to
  - Read the operands
  - Compute the result
- Parameters:
  - $B(R)$  = # of blocks (i.e., pages) for relation R
  - $T(R)$  = # of tuples in relation R

Magda Balazinska - CSE 444, Spring 2012

13

## Basic Join Algorithms

- Operator:
  - Product  $\bowtie$  Company
- Propose three algorithms for the join, assuming the tables fit in main memory:
  - Hash join
  - Nested loop join
  - Sort-merge join

Magda Balazinska - CSE 444, Spring 2012

14

## Hash Join

Hash join:  $R \bowtie S$

- Scan R, build buckets in main memory
- Then scan S and join
- Cost:  $B(R) + B(S)$

Magda Balazinska - CSE 444, Spring 2012

15

## Hash Join Example

Patient(pid, name, address)

Insurance(pid, provider, policy\_nb)

Patient  $\bowtie$  Insurance

Two tuples per page

Patient			Insurance		
1	'Bob'	'Seattle'	2	'Blue'	123
2	'Ela'	'Everett'	4	'Prem'	432
3	'Jill'	'Kent'	4	'Prem'	343
4	'Joe'	'Seattle'	3	'GrpH'	554

16

## Hash Join Example

Patient  $\bowtie$  Insurance

Some large-enough nb

Memory M = 21 pages

Showing pid only

Disk	
Patient	Insurance
1 2	2 4 6 6
3 4	4 3 1 3
9 6	2 8
8 5	8 9

17

## Hash Join Example

Step 1: Scan Patient and create hash table in memory

Memory M = 21 pages

Hash h: pid % 5

5	1	6	2	3	8	4	9
---	---	---	---	---	---	---	---

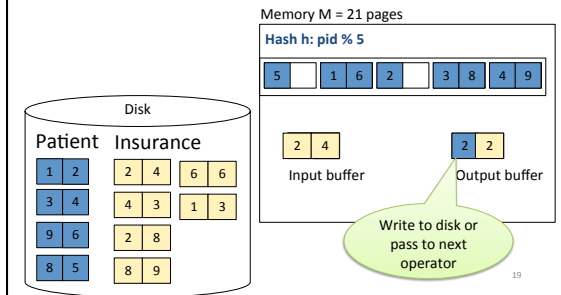
Input buffer

Disk	
Patient	Insurance
1 2	2 4 6 6
3 4	4 3 1 3
9 6	2 8
8 5	8 9

18


## Hash Join Example

Step 2: Scan Insurance and probe into hash table



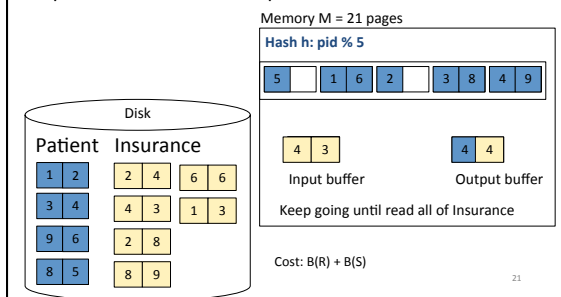
## Hash Join Example

Step 2: Scan Insurance and probe into hash table



## Hash Join Example

Step 2: Scan Insurance and probe into hash table



## Nested Loop Joins

- Tuple-based nested loop  $R \bowtie S$
- R is the outer relation, S is the inner relation

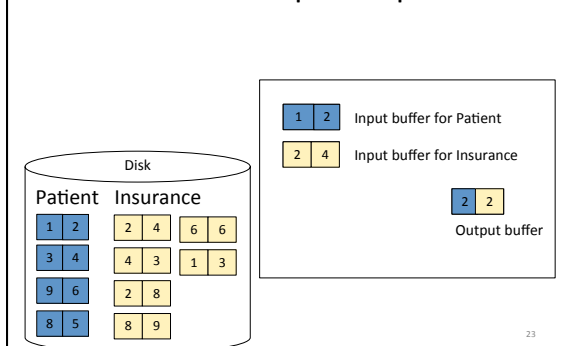
```
for each tuple r in R do
  for each tuple s in S do
    if r and s join then output (r,s)
```

- Cost:  $B(R) + B(R) B(S)$

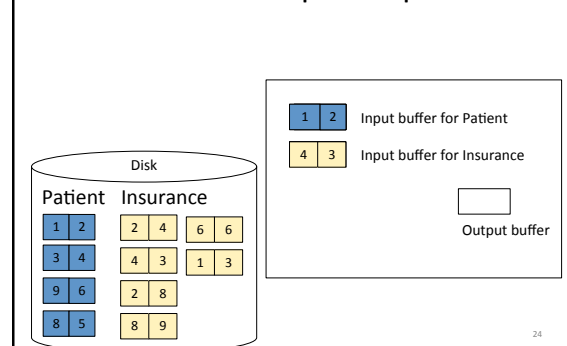
Magda Balazinska - CSE 444, Spring 2012

22

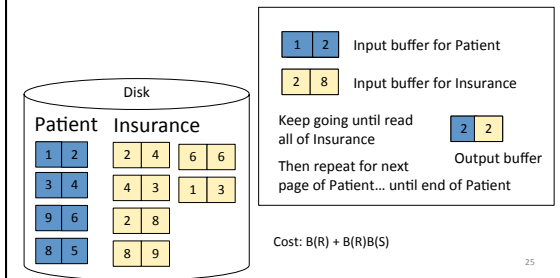
## Nested Loop Example



## Nested Loop Example



## Nested Loop Example



25

## Sort-Merge Join

Sort-merge join:  $R \bowtie S$

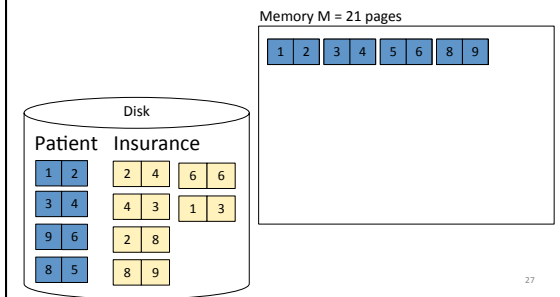
- Scan R and sort in main memory
- Scan S and sort in main memory
- Merge R and S
- Cost:  $B(R) + B(S)$

Magda Balazinska - CSE 444, Spring 2012

26

## Sort-Merge Join Example

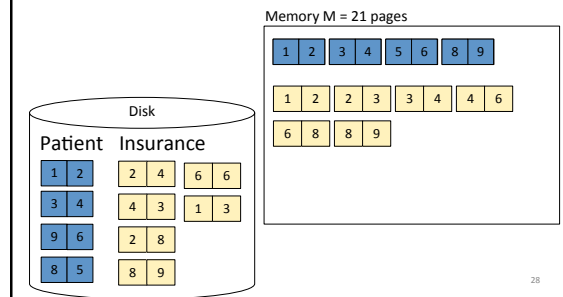
Step 1: Scan Patient and sort in memory



27

## Sort-Merge Join Example

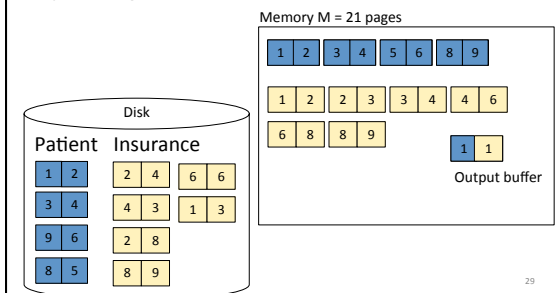
Step 2: Scan Insurance and sort in memory



28

## Sort-Merge Join Example

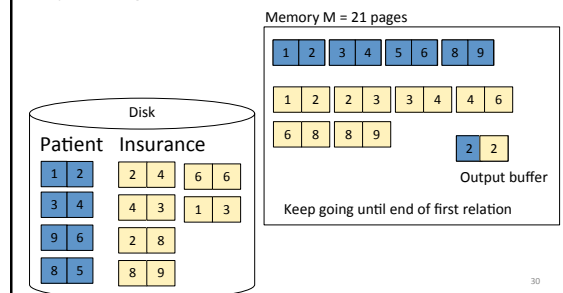
Step 3: Merge Patient and Insurance



29

## Sort-Merge Join Example

Step 3: Merge Patient and Insurance



30

## Indexes

- Special data structure for fast tuple look ups

## Example

Patient

pid	name	address
1	Bob	Seattle
2	Ela	Everett
...		

Insurance

pid	provider	Policy_nb
2	Blue	123
...		

Both tables are on disk  
How can we answer this query?

```
SELECT *
FROM Patient p, Insurance i
WHERE p.pid=i.pid AND i.provider = 'Blue'
```

Dan Suciu - CSE 344, Winter 2012

32

## Possible Query Plan 1

For all **patients**  
For all **insurance**  
If **conditions=true**  
Return **result**

Nested-loop  
join

Relations may not fit in memory.  
So may need to read courses  
many times from disk

Dan Suciu - CSE 344, Winter 2012

33

## Possible Query Plan 2

Sort **patients** on **pid**  
Sort **insurance** on **pid**  
Merge join  
Check additional conditions  
Return **results that satisfy conditions**

Sort-merge join

Sorting can take a  
long time

Dan Suciu - CSE 344, Winter 2012

34

## Possible Query Plan 3

Find and **only read** from disk insurance from **provider 'Blue'**  
For each such **policy**, find matching **patients**  
Return **results**

Can we do  
this?

Yes! But we need  
**indexes**

Dan Suciu - CSE 344, Winter 2012

35

## Index


- An **additional file**, that allows fast access to records in the data file given a search key
- The index contains (key, value) pairs:
  - The key = an attribute value (e.g., pid or name)
  - The value = a pointer to the record
- SQL statement to create index:
  - CREATE INDEX **pid\_index** ON **patient(pid)**

"key" = "search key"

Dan Suciu - CSE 344, Winter 2012

36

## Example of Index



Dan Suciu - CSE 344, Winter 2012

37

## Selections with Indexing

```
SELECT *
FROM Insurance
WHERE provider = 'Blue'
```

Without index: Read the whole of Insurance.

With index: Read only those tuples that have 'Blue' as the insurance provider, if index on **provider** exists.

## Nested Loop Joins

- Tuple-based nested loop  $R \bowtie S$
- R is the outer relation, S is the inner relation

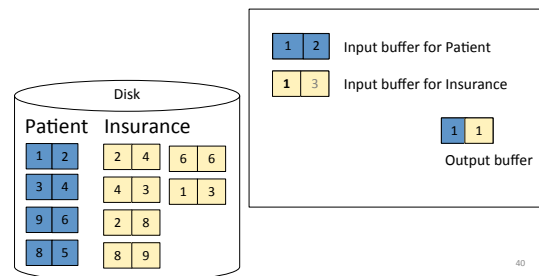
for each tuple r in R do  
     use index on S to search for key  
     then output (r,s)

- Cost:  $B(R) + B(R) * \text{constant}$

Magda Balazinska - CSE 444, Spring 2012

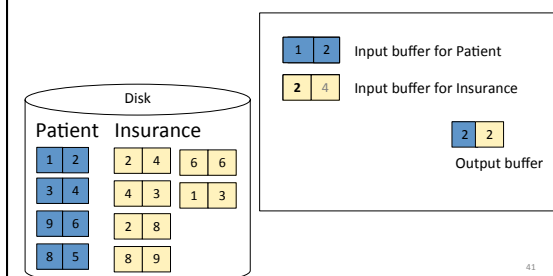
39

## Nested Loop Example



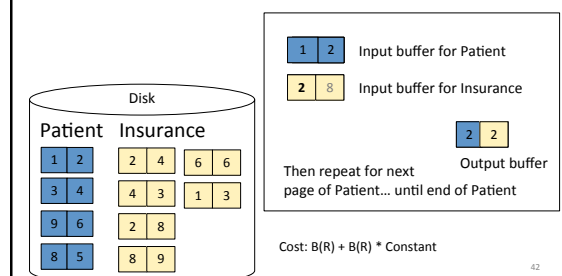
40

## Nested Loop Example



41

## Nested Loop Example



42

## Index Selection

- Depends on the query plan
- In general construct index on attributes that occur in
  - Selections
  - Join predicates
  - Grouping

## Index Selection Example

