

Introduction to Data Management CSE 344

Lectures 4 and 5: Aggregates in SQL

Magda Balazinska - CSE 344, Fall 2012

1

Announcements

- Homework 1 is due tonight!
- Homework 2 is posted (due next week)
- Quiz 1 due tonight!
- Quiz 2 due next week

Magda Balazinska - CSE 344, Fall 2012

2

Outline

- Outer joins (6.3.8)
- Aggregations (6.4.3 – 6.4.6)
- Examples, examples, examples...

Magda Balazinska - CSE 344, Fall 2012

3

Outerjoins

Product(name, category)
Purchase(prodName, store)

An "inner join":

```
SELECT Product.name, Purchase.store
FROM   Product, Purchase
WHERE  Product.name = Purchase.prodName
```

Same as:

```
SELECT Product.name, Purchase.store
FROM   Product JOIN Purchase ON
        Product.name = Purchase.prodName
```

But Products that never sold will be lost !

Magda Balazinska - CSE 344, Fall 2012

4

Outerjoins

Product(name, category)
Purchase(prodName, store)

If we want the never-sold products, need an "outerjoin":

```
SELECT Product.name, Purchase.store
FROM   Product LEFT OUTER JOIN Purchase ON
        Product.name = Purchase.prodName
```

Magda Balazinska - CSE 344, Fall 2012

5

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL

6

Outer Joins

- Left outer join:
 - Include the left tuple even if there's no match
- Right outer join:
 - Include the right tuple even if there's no match
- Full outer join:
 - Include both left and right tuples even if there's no match

Magda Balazinska - CSE 344, Fall 2012

7

Aggregation in SQL

sqlite3 lecture04

Specify a filename where the database will be stored

```
create table Purchase
(pid int primary key,
 product varchar(15),
 price float,
 quantity int,
 month varchar(15));
```

Other DBMSs have other ways of importing data

```
.import data.txt Purchase
```

Magda Balazinska - CSE 344, Fall 2012

8

Comment about SQLite

- One cannot load NULL values such that they are actually loaded as null values
- So we need to use two steps:
 - Load null values using some type of special value
 - Update the special values to actual null values

```
update Purchase set price = null where price = 'null'
```

Magda Balazinska - CSE 344, Fall 2012

9

Simple Aggregations

Five basic aggregate operations in SQL

- `select count(*) from Purchase`
- `select count(quantity) from Purchase`
- `select sum(quantity) from Purchase`
- `select avg(price) from Purchase`
- `select max(quantity) from Purchase`
- `select min(quantity) from Purchase`

Except count, all aggregations apply to a single attribute

Magda Balazinska - CSE 344, Fall 2012

10

Aggregates and NULL Values

Null values are not used in aggregates

- `insert into Purchase values(12, 'gadget', NULL, NULL, 'april')`

Let's try the following:

- `select count(*) from Purchase`
- `select count(quantity) from Purchase`
- `select sum(quantity) from Purchase`
- `select sum(quantity) from Purchase where quantity is not null;`

Magda Balazinska - CSE 344, Fall 2012

11

Counting Duplicates

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(product)
FROM Purchase
WHERE price > 4.99
```

same as Count(*)

We probably want:

```
SELECT Count(DISTINCT product)
FROM Purchase
WHERE price > 4.99
```

Magda Balazinska - CSE 344, Fall 2012

12

More Examples

```
SELECT Sum(price * quantity)
FROM Purchase
```

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```

What do they mean ?

Magda Balazinska - CSE 344, Fall 2012

13

Simple Aggregations

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'Bagel'
```

→ 90 (= 60+30)

Magda Balazinska - CSE 344, Fall 2012

14

Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantities for all sales over \$1, by product.

```
SELECT product, Sum(quantity) AS TotalSales
FROM Purchase
WHERE price > 1
GROUP BY product
```

Let's see what this means...

Magda Balazinska - CSE 344, Fall 2012

15

Grouping and Aggregation

1. Compute the **FROM** and **WHERE** clauses.
2. Group by the attributes in the **GROUPBY**
3. Compute the **SELECT** clause:
grouped attributes and aggregates.

Magda Balazinska - CSE 344, Fall 2012

16

1&2. FROM-WHERE-GROUPBY

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

WHERE price > 1

Magda Balazinska - CSE 344, Fall 2012

17

3. SELECT

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	TotalSales
Bagel	40
Banana	20

```
SELECT product, Sum(quantity) AS TotalSales
FROM Purchase
WHERE price > 1
GROUP BY product
```

18

Other Examples

Compare these two queries:

```
SELECT product, count(*)
FROM Purchase
GROUP BY product
```

```
SELECT month, count(*)
FROM Purchase
GROUP BY month
```

```
SELECT product,
sum(quantity) AS SumQuantity,
max(price) AS MaxPrice
FROM Purchase
GROUP BY product
```

What does it mean ?

Magda Balazinska - CSE 344, Fall 2012

19

Need to be Careful...

```
SELECT product, max(quantity)
FROM Purchase
GROUP BY product
```

```
SELECT product, quantity
FROM Purchase
GROUP BY product
```

Sqlite is WRONG on this query.

SQL Server correctly gives an error

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Magda Balazinska - CSE 344, Fall 2012

20

Ordering Results

```
SELECT product, sum(price*quantity) as rev
FROM purchase
GROUP BY product
ORDER BY rev desc
```

Magda Balazinska - CSE 344, Fall 2012

21

HAVING Clause

Same query as earlier, except that we consider only products that had at least 30 sales.

```
SELECT product, sum(price*quantity)
FROM Purchase
WHERE price > 1
GROUP BY product
HAVING Sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.

Magda Balazinska - CSE 344, Fall 2012

22

WHERE vs HAVING

- WHERE condition is applied to individual rows
 - The rows may or may not contribute to the aggregate
 - No aggregates allowed here
- HAVING condition is applied to the entire group
 - Entire group is returned, or not at all
 - May use aggregate functions in the group

Magda Balazinska - CSE 344, Fall 2012

23

Aggregates and Joins

```
create table Product (pid int primary key, pname
varchar(15), manufacturer varchar(15));
```

```
insert into product values(1, 'bagel', 'Sunshine Co. ');
insert into product values(2, 'banana', 'BusyHands');
insert into product values(3, 'gizmo', 'GizmoWorks');
insert into product values(4, 'gadget', 'BusyHands');
insert into product values(5, 'powerGizmo', 'PowerWorks');
```

Magda Balazinska - CSE 344, Fall 2012

24

Aggregate + Join Example

```
SELECT x.manufacturer, count(*)
FROM Product x, Purchase y
WHERE x.pname = y.product
GROUP BY x.manufacturer
```

What do these query mean?

```
SELECT x.manufacturer, y.month, count(*)
FROM Product x, Purchase y
WHERE x.pname = y.product
GROUP BY x.manufacturer, y.month
```

Magda Balazinska - CSE 344, Fall 2012

25

General form of Grouping and Aggregation

```
SELECT S
FROM R1, ..., Rn
WHERE C1
GROUP BY a1, ..., ak
HAVING C2
```

Why?

S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES
 C1 = is any condition on the attributes in R_1, \dots, R_n
 C2 = is any condition on aggregate expressions and on attributes a_1, \dots, a_k

Magda Balazinska - CSE 344, Fall 2012

26

Semantics of SQL With Group-By

```
SELECT S
FROM R1, ..., Rn
WHERE C1
GROUP BY a1, ..., ak
HAVING C2
```

Evaluation steps:

1. Evaluate FROM-WHERE using Nested Loop Semantics
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

Magda Balazinska - CSE 344, Fall 2012

27

Empty Groups

- In the result of a group by query, there is one row per group in the result
- No group can be empty!
- In particular, count(*) is never 0

What if there are no purchases for a manufacturer

```
SELECT x.manufacturer, count(*)
FROM Product x, Purchase y
WHERE x.pname = y.product
GROUP BY x.manufacturer
```

Magda Balazinska - CSE 344, Fall 2012

28

Empty Groups: Example

```
SELECT product, count(*)
FROM purchase
GROUP BY product
```

5 groups in our example dataset

```
SELECT product, count(*)
FROM purchase
WHERE price > 2.0
GROUP BY product
```

3 groups in our example dataset

Magda Balazinska - CSE 344, Fall 2012

29

Empty Group Problem

```
SELECT x.manufacturer, count(*)
FROM Product x, Purchase y
WHERE x.pname = y.product
GROUP BY x.manufacturer
```

What if there are no purchases for a manufacturer

Magda Balazinska - CSE 344, Fall 2012

30

Empty Group Solution: Outer Join

```
SELECT x.manufacturer, count(y.pid)
FROM Product x LEFT OUTER JOIN Purchase y
ON x.pname = y.product
GROUP BY x.manufacturer
```