

Introduction to Database Systems CSE 444

Lecture 6: Basic Database Tuning

Magda Balazinska - CSE 344, Fall 2012

1

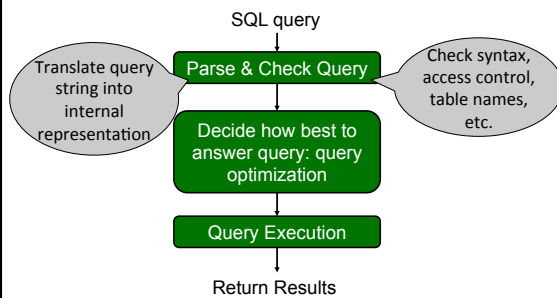
Where We Are

- We learned importance and benefits of DBMSs
- We learned how to use a DBMS
 - How to specify what our data will look like: schema
 - How to load data into the DBMS
 - How to ask simple select-project-join-agg. queries
- Today: how to get queries to run faster

Magda Balazinska - CSE 344, Fall 2012

2

Query Evaluation Steps



Magda Balazinska - CSE 344, Fall 2012

3

Example

Student

ID	fName	lName
195428	Tom	Hanks
645947	Amy	Hanks
...		

Courses

studentID	courseID
195428	344
...	

Both tables are on disk
How can we answer this query?


```

SELECT *
FROM Student S, Courses C
WHERE S.ID=C.studentID AND C.courseID >= 300
  
```

Magda Balazinska - CSE 344, Fall 2012

4


Possible Query Plan 1



Magda Balazinska - CSE 344, Fall 2012

5

Possible Query Plan 2



Magda Balazinska - CSE 344, Fall 2012

6

Possible Query Plan 3

Create a hash-table of students on ID
Read courses and probe hash table
If match found, check additional conditions
Return results that satisfy the conditions

Hash-join

Still have to read
entire relations
from disk!

Hash table
may not fit in
memory

Magda Balazinska - CSE 344, Fall 2012

7

Possible Query Plan 4

Find and only read from disk courses with courseID ≥ 300
For each such course, find matching students
Return results

Can we do
this?

Yes! But we
need **indexes**

Magda Balazinska - CSE 344, Fall 2012

8

Data Storage

- DBMSs store data in **files**
- Most common organization is row-wise storage
- On disk, a file is split into blocks
- Each block contains a set of tuples

10	...
20	
30	
40	
50	
60	
70	
80	

In the example, we have 4 blocks with 2 tuples each

Magda Balazinska - CSE 344, Fall 2012

9

Database File Types

The data file can be one of:

- Heap file**
 - Unsorted
- Sequential file**
 - Sorted according to some attribute(s) called **key**

"key" here means something else than "primary key"
Example: ID is primary key for students
But can sort students on last name

Magda Balazinska - CSE 344, Fall 2012

10

Index


- An **additional file**, that allows fast access to records in the data file given a search key
- The index contains (key, value) pairs:
 - The key = an attribute value (e.g., student ID or name)
 - The value = a pointer to the record

"key" = "search key"

Magda Balazinska - CSE 344, Fall 2012

11


Example of Index



Magda Balazinska - CSE 344, Fall 2012

12

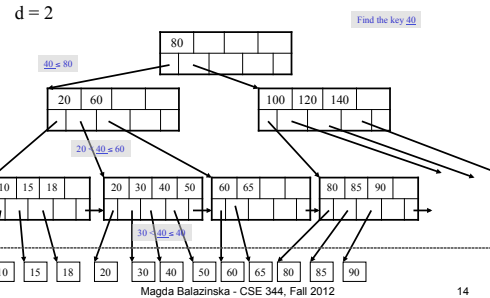
Hash-Based Index by Example



Magda Balazinska - CSE 344, Fall 2012

13


B+ Tree Index by Example



Magda Balazinska - CSE 344, Fall 2012

14

Clustered vs Unclustered



Magda Balazinska - CSE 344, Fall 2012

15

Index Classification

- **Clustered/unclustered**
 - Clustered = records close in index are close in data
 - Option 1: Data inside data file is sorted on disk
 - Option 2: Store data directly inside the index (no separate files)
 - Unclustered = records close in index may be far in data
- **Primary/secondary**
 - Meaning 1:
 - Primary = is over attributes that include the primary key
 - Secondary = otherwise
 - Meaning 2: means the same as clustered/unclustered
- **Organization:** B+ tree or Hash table

Magda Balazinska - CSE 344, Fall 2012

16

Indexes in SQL

```
CREATE TABLE V(M int, N varchar(20), P int);
```

```
CREATE INDEX V1 ON V(N)
```

```
CREATE INDEX V2 ON V(P, M)
```

```
CREATE INDEX V3 ON V(M, N)
```

```
CREATE UNIQUE INDEX V4 ON V(N)
```

```
CREATE CLUSTERED INDEX V5 ON V(N)
```

OK in SQL Server
but not supported
in SQLite

Magda Balazinska - CSE 344, Fall 2012

17

The Index Selection Problem

- Given a database schema (tables, attributes)
- Given a "query workload":
 - Workload = a set of (query, frequency) pairs
 - The queries may be both SELECT and updates
 - Frequency = either a count, or a percentage
- Select a set of indexes that optimizes the workload

In general this is a very hard problem

Magda Balazinska - CSE 344, Fall 2012

18

Index Selection: Which Search Key

- Make some attribute K a search key if the WHERE clause contains:
 - An exact match on K
 - A range predicate on K
 - A join on K

Magda Balazinska - CSE 344, Fall 2012

19

The Index Selection Problem 1

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

What indexes ?

Magda Balazinska - CSE 344, Fall 2012

20

The Index Selection Problem 1

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

A: $V(N)$ and $V(P)$ (hash tables or B-trees)

Magda Balazinska - CSE 344, Fall 2012

21

The Index Selection Problem 2

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N>? and N<?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

100000 queries:

```
INSERT INTO V
VALUES (?, ?, ?)
```

What indexes ?

Magda Balazinska - CSE 344, Fall 2012

22

The Index Selection Problem 2

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N>? and N<?
```

100 queries:

```
SELECT *
FROM V
WHERE P=?
```

100000 queries:

```
INSERT INTO V
VALUES (?, ?, ?)
```

A: definitely $V(N)$ (must B-tree); unsure about $V(P)$

Magda Balazinska - CSE 344, Fall 2012

23

The Index Selection Problem 3

$V(M, N, P);$

Your workload is this

100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

1000000 queries:

```
SELECT *
FROM V
WHERE N=? and P>?
```

100000 queries:

```
INSERT INTO V
VALUES (?, ?, ?)
```

What indexes ?

Magda Balazinska - CSE 344, Fall 2012

24

The Index Selection Problem 3

V(M, N, P);

Your workload is this

100000 queries: 1000000 queries: 100000 queries:

```
SELECT *
FROM V
WHERE N=?
```

```
SELECT *
FROM V
WHERE N=? and P>?
```

```
INSERT INTO V
VALUES (?, ?, ?)
```

A: V(N, P)

Magda Balazinska - CSE 344, Fall 2012

25

The Index Selection Problem 4

V(M, N, P);

Your workload is this

1000 queries:

100000 queries:

```
SELECT *
FROM V
WHERE N>? and N<?
```

```
SELECT *
FROM V
WHERE P>? and P<?
```

What indexes ?

Magda Balazinska - CSE 344, Fall 2012

26

The Index Selection Problem 4

V(M, N, P);

Your workload is this

1000 queries:

100000 queries:

```
SELECT *
FROM V
WHERE N>? and N<?
```

```
SELECT *
FROM V
WHERE P>? and P<?
```

A: V(N) secondary, V(P) primary index

Magda Balazinska - CSE 344, Fall 2012

27

The Index Selection Problem

- **SQL Server**
 - Automatically, thanks to *AutoAdmin* project
 - Much acclaimed successful research project from mid 90's, similar ideas adopted by the other major vendors
 - But can also do this manually
- **SQLite**
 - You will do it manually, part of homework 2

Magda Balazinska - CSE 344, Fall 2012

28

Basic Index Selection Guidelines

- Consider queries in workload in order of importance
- Consider relations accessed by query
 - No point indexing other relations
- Look at WHERE clause for possible search key
- Try to choose indexes that speed-up multiple queries
- And then consider the following...

Magda Balazinska - CSE 344, Fall 2012

29

Index Selection: Multi-attribute Keys

Consider creating a multi-attribute key on K1, K2, ... if

- WHERE clause has matches on K1, K2, ...
 - But also consider separate indexes
- SELECT clause contains only K1, K2, ..
 - A *covering index* is one that can be used exclusively to answer a query, e.g. index R(K1,K2) covers the query:

```
SELECT K2 FROM R WHERE K1=55
```

Magda Balazinska - CSE 344, Fall 2012

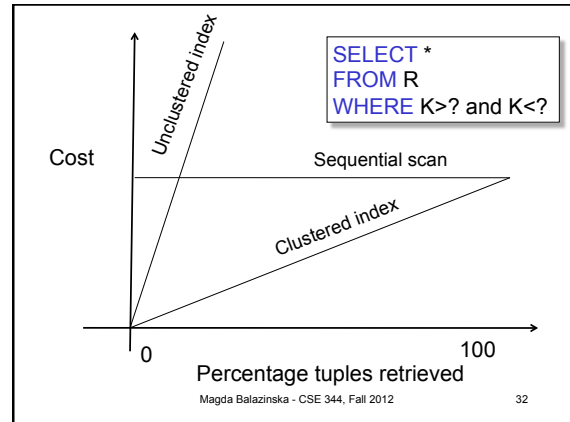
30

To Cluster or Not

- Range queries benefit mostly from clustering
- Covering indexes do *not* need to be clustered: they work equally well unclustered

Magda Balazinska - CSE 344, Fall 2012

31



Magda Balazinska - CSE 344, Fall 2012

32

Hash Table v.s. B+ tree

- Rule 1: always use a B+ tree ☺
- Rule 2: use a Hash table on K when:
 - There is a very important selection query on equality (WHERE K=?), and no range queries
 - You know that the optimizer uses a nested loop join where K is the join attribute of the inner relation

Magda Balazinska - CSE 344, Fall 2012

33

Balance Queries v.s. Updates

- Indexes speed up queries
 - SELECT FROM WHERE
- But they usually slow down updates:
 - INSERT, DELETE, UPDATE
 - However some updates benefit from indexes

```
UPDATE R
SET A = 7
WHERE K=55
```

Magda Balazinska - CSE 344, Fall 2012

34

Tools for Index Selection

- SQL Server 2000 Index Tuning Wizard
- DB2 Index Advisor
- How they work:
 - They walk through a large number of configurations, compute their costs, and choose the configuration with minimum cost

Magda Balazinska - CSE 344, Fall 2012

35

The Database Tuning Problem

- We are given a workload description
 - List of queries and their frequencies
 - List of updates and their frequencies
 - Performance goals for each type of query
- Perform *physical database design*
 - Choose indexes
 - Other tunings are also possible

Magda Balazinska - CSE 344, Fall 2012

36