

# Introduction to Data Management CSE 344

## Lecture 7: Nested Queries in SQL

Magda Balazinska - CSE 344, Fall 2012

1

## Lecture Goals

- Today we will learn how to write more powerful SQL queries
- They are needed in Homework 3
- Reminder: Book chapters associated with lectures are listed on the calendar page of the course website

Magda Balazinska - CSE 344, Fall 2012

2

## Subqueries

- A subquery is a SQL query nested inside a larger query
- Such inner-outer queries are called nested queries
- A subquery may occur in:
  - A SELECT clause
  - A FROM clause
  - A WHERE clause
- Rule of thumb: avoid writing nested queries when possible; keep in mind that sometimes it's impossible

Magda Balazinska - CSE 344, Fall 2012

3

## 1. Subqueries in SELECT

Product (pname, price, cid)  
Company(cid, cname, city)

For each product return the city where it is manufactured

```
SELECT X.pname, (SELECT Y.city
                  FROM Company Y
                  WHERE Y.cid=X.cid) as City
FROM Product X
```

"correlated subquery"

What happens if the subquery returns more than one city ?

We get a runtime error  
(SQLite simply ignores the extra values)

Magda Balazinska - CSE 344, Fall 2012

4

## 1. Subqueries in SELECT

Product (pname, price, cid)  
Company(cid, cname, city)

Whenever possible, don't use a nested queries:

```
SELECT X.pname, (SELECT Y.city
                  FROM Company Y
                  WHERE Y.cid=X.cid) as City
FROM Product X
```

=

```
SELECT X.pname, Y.city
FROM Product X, Company Y
WHERE X.cid=Y.cid
```

We have  
"unnested"  
the query

5

## 1. Subqueries in SELECT

Product (pname, price, cid)  
Company(cid, cname, city)

Compute the number of products made by each company

```
SELECT DISTINCT C.cname, (SELECT count(*)
                           FROM Product P
                           WHERE P.cid=C.cid)
FROM Company C
```

Better: we can  
unnest by using  
a GROUP BY

```
SELECT C.cname, count(*)
FROM Company C, Product P
WHERE C.cid=P.cid
GROUP BY C.cname
```

Magda Balazinska - CSE 344, Fall 2012

6

## 1. Subqueries in SELECT

Are these really equivalent?

```
SELECT DISTINCT C.cname, (SELECT count(*)  
FROM Product P  
WHERE P.cid=C.cid)  
FROM Company C
```

```
SELECT C.cname, count(*)  
FROM Company C, Product P  
WHERE C.cid=P.cid  
GROUP BY C.cname
```

No! Different results if a company has no products

```
SELECT C.cname, count(pname)  
FROM Company C LEFT OUTER JOIN Product P  
ON C.cid=P.cid  
GROUP BY C.cname
```

Magda Balazinska - CSE 344, Fall 2012

7

## 2. Subqueries in FROM

Product (pname, price, cid)  
Company(cid, cname, city)

Find all products whose prices is > 20 and < 500

```
SELECT X.pname  
FROM (SELECT * FROM Product AS Y WHERE price > 20) as X  
WHERE X.price < 500
```

Unnest this query !

Magda Balazinska - CSE 344, Fall 2012

8

## 2. Subqueries in FROM

- At the end of the lecture we will see that sometimes we really need a subquery and one option will be to put it in the FROM clause (see "finding witnesses").

Magda Balazinska - CSE 344, Fall 2012

9

## 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Existential quantifiers

Find all companies that make some products with price < 200

Using EXISTS:

```
SELECT DISTINCT C.cname  
FROM Company C  
WHERE EXISTS (SELECT *  
FROM Product P  
WHERE C.cid = P.cid and P.price < 200)
```

Magda Balazinska - CSE 344, Fall 2012

10

## 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Existential quantifiers

Find all companies that make some products with price < 200

Using IN

```
SELECT DISTINCT C.cname  
FROM Company C  
WHERE C.cid IN (SELECT P.cid  
FROM Product P  
WHERE P.price < 200)
```

Magda Balazinska - CSE 344, Fall 2012

11

## 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Existential quantifiers

Find all companies that make some products with price < 200

Using ANY:

```
SELECT DISTINCT C.cname  
FROM Company C  
WHERE 200 > ANY (SELECT price  
FROM Product P  
WHERE P.cid = C.cid)
```

Magda Balazinska - CSE 344, Fall 2012

12

### 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Existential quantifiers

Find all companies that make some products with price < 200

Now let's unnest it:

```
SELECT DISTINCT C.cname
FROM   Company C, Product P
WHERE  C.cid= P.cid and P.price < 200
```

Existential quantifiers are easy ! ☺

13

### 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Universal quantifiers

Find all companies that make only products with price < 200

same as:

Find all companies whose products all have price < 200

Universal quantifiers are hard ! ☹

Magda Balazinska - CSE 344, Fall 2012

14

### 3. Subqueries in WHERE

1. Find *the other* companies: i.e. s.t. some product ≥ 200

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  C.cid IN (SELECT P.cid
                FROM Product P
                WHERE P.price >= 200)
```

2. Find all companies s.t. all their products have price < 200

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  C.cid NOT IN (SELECT P.cid
                    FROM Product P
                    WHERE P.price >= 200)
```

15

### 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Universal quantifiers

Find all companies that make only products with price < 200

Using EXISTS:

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  NOT EXISTS (SELECT *
                  FROM Product P
                  WHERE P.cid = C.cid and P.price >= 200)
```

Magda Balazinska - CSE 344, Fall 2012

16

### 3. Subqueries in WHERE

Product (pname, price, cid)  
Company(cid, cname, city)

Universal quantifiers

Find all companies that make only products with price < 200

Using ALL:

```
SELECT DISTINCT C.cname
FROM   Company C
WHERE  200 > ALL (SELECT price
                 FROM Product P
                 WHERE P.cid = C.cid)
```

Magda Balazinska - CSE 344, Fall 2012

17

### Question for Database Fans and their Friends

- Can we unnest the *universal quantifier* query ?

Magda Balazinska - CSE 344, Fall 2012

18

## Monotone Queries

- A query Q is **monotone** if:
  - Whenever we add tuples to one or more of the tables...
  - ... the answer to the query cannot contain fewer tuples
- Fact:** all unnested queries are monotone
  - Proof: using the "nested for loops" semantics
- Fact:** Query with universal quantifier is not monotone
- Consequence:** we cannot unnest a query with a universal quantifier

Magda Balazinska - CSE 344, Fall 2012

19

## Queries that must be nested

- Queries with universal quantifiers or with negation
- The drinkers-bars-beers example next
- This is a famous example from textbook on databases by Ullman

Magda Balazinska - CSE 344, Fall 2012

20

## The drinkers-bars-beers example

Likes(drinker, beer)  
Frequents(drinker, bar)  
Serves(bar, beer)

Challenge: write these in SQL

Find drinkers that frequent some bar that serves some beer they like.

$x: \exists y. \exists z. \text{Frequents}(x, y) \wedge \text{Serves}(y, z) \wedge \text{Likes}(x, z)$

Find drinkers that frequent only bars that serves some beer they like.

$x: \forall y. \text{Frequents}(x, y) \Rightarrow (\exists z. \text{Serves}(y, z) \wedge \text{Likes}(x, z))$

Find drinkers that frequent some bar that serves only beers they like.

$x: \exists y. \text{Frequents}(x, y) \wedge \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$

Find drinkers that frequent only bars that serves only beer they like.

$x: \forall y. \text{Frequents}(x, y) \Rightarrow \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$  <sup>21</sup>

## GROUP BY v.s. Nested Queries

```
SELECT product, Sum(quantity) AS TotalSales
FROM Purchase
WHERE price > 1
GROUP BY product
```

```
SELECT DISTINCT x.product, (SELECT Sum(y.quantity)
FROM Purchase y
WHERE x.product = y.product
AND price > 1)
AS TotalSales
FROM Purchase x
WHERE price > 1
```

Why twice ? <sup>22</sup>

## Unnesting Aggregates

Product ( pname, price, cid)  
Company(cid, cname, city)

Find the number of companies in each city

```
SELECT DISTINCT city, (SELECT count(*)
FROM Company Y
WHERE X.city = Y.city)
FROM Company X
```

```
SELECT city, count(*)
FROM Company
GROUP BY city
```

Equivalent queries

Note: no need for DISTINCT (DISTINCT *is the same* as GROUP BY) <sup>23</sup>

Magda Balazinska - CSE 344, Fall 2012

## Unnesting Aggregates

Product ( pname, price, cid)  
Company(cid, cname, city)

Find the number of products made in each city

```
SELECT DISTINCT X.city, (SELECT count(*)
FROM Product Y, Company Z
WHERE Z.cid=Y.cid
AND Z.city = X.city)
FROM Company X
```

```
SELECT X.city, count(*)
FROM Company X, Product Y
WHERE X.cid=Y.cid
GROUP BY X.city
```

They are NOT equivalent ! (WHY?)

24

## More Unnesting

Author(login,name)  
Wrote(login,url)

- Find authors who wrote  $\geq 10$  documents:
- Attempt 1: with nested queries

```
SELECT DISTINCT Author.name
FROM Author
WHERE (SELECT count(Wrote.url)
FROM Wrote
WHERE Author.login=Wrote.login)
> 10
```

This is  
SQL by  
a novice

Magda Balazinska - CSE 344, Fall 2012

25

## More Unnesting

- Find all authors who wrote at least 10 documents:
- Attempt 2: SQL style (with GROUP BY)

```
SELECT Author.name
FROM Author, Wrote
WHERE Author.login=Wrote.login
GROUP BY Author.name
HAVING count(wrote.url) > 10
```

This is  
SQL by  
an expert

Magda Balazinska - CSE 344, Fall 2012

26

## Finding Witnesses

Product ( pname, price, cid)  
Company(cid, cname, city)

For each city, find the most expensive product made in that city

Magda Balazinska - CSE 344, Fall 2012

27

## Finding Witnesses

Product ( pname, price, cid)  
Company(cid, cname, city)

For each city, find the most expensive product made in that city  
Finding the maximum price is easy...

```
SELECT x.city, max(y.price)
FROM Company x, Product y
WHERE x.cid = y.cid
GROUP BY x.city;
```

But we need the *witnesses*, i.e. the products with max price

Magda Balazinska - CSE 344, Fall 2012

28

## Finding Witnesses

And another one:

```
SELECT u.city, v.pname, v.price
FROM Company u, Product v
WHERE u.cid = v.cid
and v.price >= ALL (SELECT y.price
FROM Company x, Product y
WHERE u.city=x.city
and x.cid=y.cid);
```

Magda Balazinska - CSE 344, Fall 2012

29

## Finding Witnesses

To find the witnesses, compute the maximum price in a subquery

```
SELECT DISTINCT u.city, v.pname, v.price
FROM Company u, Product v,
(SELECT x.city, max(y.price) as maxprice
FROM Company x, Product y
WHERE x.cid = y.cid
GROUP BY x.city) w
WHERE u.cid = v.cid
and u.city = w.city
and v.price=w.maxprice;
```

Magda Balazinska - CSE 344, Fall 2012

30

## Finding Witnesses

There is a more concise solution here:

```
SELECT u.city, v.pname, v.price
FROM Company u, Product v, Company x, Product y
WHERE u.cid = v.cid and u.city = x.city and x.cid = y.cid
GROUP BY u.city, v.pname, v.price
HAVING v.price = max(y.price);
```