

Introduction to Data Management CSE 344

Lecture 11: Relational Calculus

Magda Balazinska - CSE 344, Fall 2012

1

But First...

A few additional datalog examples

Friend(name1, name2)

Enemy(name1, name2)

Find Joe's friends, and Joe's friends of friends.

```
A(x) :- Friend('Joe', x)
A(x) :- Friend('Joe', z), Friend(z, x)
```

Magda Balazinska - CSE 344, Fall 2012

2

Datalog Example 2

Friend(name1, name2)

Enemy(name1, name2)

Find all of Joe's friends who do not have any friends except for Joe:

```
JoeFriends(x) :- Friend('Joe', x)
NonAns(x) :- Friend(x, y), y != 'Joe'
A(x) :- JoeFriends(x) NOT NonAns(x)
```

Magda Balazinska - CSE 344, Fall 2012

3

```
Friend(name1, name2)
Enemy(name1, name2)
```

Datalog Example 3

Find all x such that all their enemies' enemies are their friends

- Assume that if someone doesn't have any enemies nor friends, we also want them in the answer

```
Everyone(x) :- Friend(x, y)
Everyone(x) :- Friend(y, x)
Everyone(x) :- Enemy(x, y)
Everyone(x) :- Enemy(y, x)
NonAns(x) :- Enemy(x, y), Enemy(y, z) NOT Friend(x, z)
A(x) :- Everyone(x) NOT NonAns(x)
```

```
Friend(name1, name2)
Enemy(name1, name2)
```

Datalog Example 4

Find all x having some friend all of whose enemies are x's enemies.

```
Everyone(x) :- Friend(x, y)
NonAns(x) :- Friend(x, y) Enemy(y, z) NOT Enemy(x, z)
A(x) :- Everyone(x) NOT NonAns(x)
```

Magda Balazinska - CSE 344, Fall 2012

5

Why Did We Learn Datalog?

- Simple, logic language, based on rules
- Can be extended to recursion BUT beyond 344
- Equivalences
 - Datalog can be translated to SQL (practice at home !)
 - Can also translate back and forth between datalog and relational algebra (see last lecture)
 - Bottom line: relational algebra, non-recursive datalog with negation, and relational calculus all have the same expressive power!

Magda Balazinska - CSE 344, Fall 2012

6

Why Did We Learn Datalog?

Datalog, RA, and RC are of fundamental importance in DBMSs because

1. Sufficiently expressive to be useful in practice yet
2. Sufficiently simple to be efficiently implementable

Magda Balazinska - CSE 344, Fall 2012

7

Relational Calculus

- Aka predicate calculus or first order logic
- The most expressive formalism for queries: easy to write complex queries
- TRC = Tuple RC = named perspective
 - We study this one only
- DRC = Domain RC = unnamed perspective
 - Good to know that it also exists

Magda Balazinska - CSE 344, Fall 2012

8

Relational Calculus

Relational predicate P is a formula given by this grammar:

$P ::= \text{atom} \mid P \wedge P \mid P \vee P \mid P \Rightarrow P \mid \text{not}(P) \mid \forall x.P \mid \exists x.P$

Query Q :

$Q(x_1, \dots, x_k) = P$

Example: find the first/last names of actors who acted in 1940

$Q(f,l) = \exists x. \exists y. \exists z. (\text{Actor}(z,f,l) \wedge \text{Casts}(z,x) \wedge \text{Movie}(x,y,1940))$

What does this query return ?

$Q(f,l) = \exists z. (\text{Actor}(z,f,l) \wedge \forall x. (\text{Casts}(z,x) \Rightarrow \exists y. \text{Movie}(x,y,1940)))$ 9

Important Observation

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Find all bars that serve all beers that Fred likes

$A(x) = \forall y. \text{Likes}(\text{"Fred"}, y) \Rightarrow \text{Serves}(x,y)$

- Note: $P \Rightarrow Q$ (read P implies Q) is the same as $(\text{not } P) \text{ OR } Q$
In this query: If Fred likes a beer the bar must serve it ($P \Rightarrow Q$)
In other words: Either Fred does not like the beer ($\text{not } P$) OR the bar serves that beer (Q).

$A(x) = \forall y. \text{not}(\text{Likes}(\text{"Fred"}, y)) \text{ OR } \text{Serves}(x,y)$

Magda Balazinska - CSE 344, Fall 2012

10

More Examples

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Find drinkers that frequent some bar that serves some beer they like.

$Q(x) = \exists y. \exists z. \text{Frequents}(x, y) \wedge \text{Serves}(y,z) \wedge \text{Likes}(x,z)$

Magda Balazinska - CSE 344, Fall 2012

11

More Examples

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Find drinkers that frequent some bar that serves some beer they like.

$Q(x) = \exists y. \exists z. \text{Frequents}(x, y) \wedge \text{Serves}(y,z) \wedge \text{Likes}(x,z)$

Find drinkers that frequent only bars that serves some beer they like.

$Q(x) = \forall y. \text{Frequents}(x, y) \Rightarrow (\exists z. \text{Serves}(y,z) \wedge \text{Likes}(x,z))$

Magda Balazinska - CSE 344, Fall 2012

12

More Examples

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Find drinkers that frequent some bar that serves some beer they like.

$Q(x) = \exists y. \exists z. \text{Frequents}(x, y) \wedge \text{Serves}(y, z) \wedge \text{Likes}(x, z)$

Find drinkers that frequent only bars that serves some beer they like.

$Q(x) = \forall y. \text{Frequents}(x, y) \Rightarrow (\exists z. \text{Serves}(y, z) \wedge \text{Likes}(x, z))$

Find drinkers that frequent some bar that serves only beers they like.

$Q(x) = \exists y. \text{Frequents}(x, y) \wedge \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$

Magda Balazinska - CSE 344, Fall 2012

13

More Examples

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Find drinkers that frequent some bar that serves some beer they like.

$Q(x) = \exists y. \exists z. \text{Frequents}(x, y) \wedge \text{Serves}(y, z) \wedge \text{Likes}(x, z)$

Find drinkers that frequent only bars that serves some beer they like.

$Q(x) = \forall y. \text{Frequents}(x, y) \Rightarrow (\exists z. \text{Serves}(y, z) \wedge \text{Likes}(x, z))$

Find drinkers that frequent some bar that serves only beers they like.

$Q(x) = \exists y. \text{Frequents}(x, y) \wedge \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$

Find drinkers that frequent only bars that serves only beer they like.

$Q(x) = \forall y. \text{Frequents}(x, y) \Rightarrow \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$

Relational Calculus

How to write a complex SQL query:

- Write it in RC
- Translate RC to datalog (see next)
- Translate datalog to SQL

Take shortcuts when you know what you're doing

Magda Balazinska - CSE 344, Fall 2012

16

Domain Independent Relational Calculus

- As in datalog, one can write "unsafe" RC queries; they are also called domain dependent

$A(x) = \text{not Likes}(\text{"Fred"}, x)$
 $A(x, y) = \text{Likes}(\text{"Fred"}, x) \text{ OR Serves}(\text{"Bar"}, y)$

- Lesson: make sure your RC queries are domain independent

Magda Balazinska - CSE 344, Fall 2012

15

From RC to Non-recursive Datalog w/ negation

Query: Find drinkers that like some beer so much that they frequent all bars that serve it

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \forall z. (\text{Serves}(z, y) \Rightarrow \text{Frequents}(x, z))$

Magda Balazinska - CSE 344, Fall 2012

17

From RC to Non-recursive Datalog w/ negation

Query: Find drinkers that like some beer so much that they frequent all bars that serve it

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \forall z. (\text{Serves}(z, y) \Rightarrow \text{Frequents}(x, z))$

Step 1: Replace \forall with \exists using de Morgan's Laws

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \neg \exists z. (\text{Serves}(z, y) \wedge \neg \text{Frequents}(x, z))$

$\forall x P(x)$ same as $\neg \exists x \neg P(x)$

$\neg(\neg P \vee Q)$ same as $P \wedge \neg Q$

Magda Balazinska - CSE 344, Fall 2012

18

From RC to Non-recursive Datalog w/ negation

Query: Find drinkers that like some beer so much that they frequent all bars that serve it

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \forall z. (\text{Serves}(z, y) \Rightarrow \text{Frequents}(x, z))$

Step 1: Replace \forall with \exists using de Morgan's Laws

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \neg \exists z. (\text{Serves}(z, y) \wedge \neg \text{Frequents}(x, z))$

Step 2: Make all *subqueries* domain independent

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \neg \exists z. (\text{Likes}(x, y) \wedge \text{Serves}(z, y) \wedge \neg \text{Frequents}(x, z))$

Magda Balazinska - CSE 344, Fall 2012

19

From RC to Non-recursive Datalog w/ negation

$Q(x) = \exists y. \text{Likes}(x, y) \wedge \neg \exists z. (\text{Likes}(x, y) \wedge \text{Serves}(z, y) \wedge \neg \text{Frequents}(x, z))$

$H(x, y)$

Step 3: Create a datalog rule for each subexpression;
(shortcut: only for "important" subexpressions)

$H(x, y) \text{ :- Likes}(x, y), \text{Serves}(y, z), \text{not Frequents}(x, z)$

$Q(x) \text{ :- Likes}(x, y), \text{not } H(x, y)$

Magda Balazinska - CSE 344, Fall 2012

20

From RC to Non-recursive Datalog w/ negation

$H(x, y) \text{ :- Likes}(x, y), \text{Serves}(y, z), \text{not Frequents}(x, z)$

$Q(x) \text{ :- Likes}(x, y), \text{not } H(x, y)$

Step 4: Write it in SQL

SELECT DISTINCT L.drinker **FROM** Likes L
WHERE not exists
(**SELECT** * **FROM** Likes L2, Serves S
WHERE L2.drinker=L.drinker **and** L2.beer=L.beer
and L2.beer=S.beer
and not exists (**SELECT** * **FROM** Frequents F
WHERE F.drinker=L2.drinker
and F.bar=S.bar))

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

21

From RC to Non-recursive Datalog w/ negation

$H(x, y) \text{ :- Likes}(x, y), \text{Serves}(y, z), \text{not Frequents}(x, z)$

$Q(x) \text{ :- Likes}(x, y), \text{not } H(x, y)$

Unsafe rule

Improve the SQL query by using an unsafe datalog rule

SELECT DISTINCT L.drinker **FROM** Likes L
WHERE not exists
(**SELECT** * **FROM** Serves S
WHERE L.beer=S.beer
and not exists (**SELECT** * **FROM** Frequents F
WHERE F.drinker=L.drinker
and F.bar=S.bar))

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Magda Balazinska - CSE 344, Fall 2012

22

Summary of Translation

- RC \rightarrow recursion-free datalog w/ negation
 - Subtle: as we saw; more details in the paper
- Recursion-free datalog w/ negation \rightarrow RA
- RA \rightarrow RC

Theorem: RA, non-recursive datalog w/ negation, and RC, express exactly the same sets of queries:
RELATIONAL QUERIES

Magda Balazinska - CSE 344, Fall 2012

23