

Introduction to Data Management CSE 344

Lecture 17: Views

Magda Balazinska - CSE 344, Fall 2012

1

What is a View?

A view is a relation defined by a query

Purchase(customer, product, store) StorePrice(store, price)
Product(pname, price)

```
CREATE VIEW StorePrice AS
SELECT x.store, y.price
FROM Purchase x, Product y
WHERE x.pid = y.pid
```

This is like a new table
StorePrice(store, price)

Magda Balazinska - CSE 344, Fall 2012

2

Customer(cid, name, city) StorePrice(store, price)
Purchase(customer, product, store)
Product(pname, price)

How to Use a View?

- A "high end" store is a store that sold some product over 1000. For each customer, find all the high end stores that they visit. Return a set of (customer-name, high-end-store) pairs.

```
SELECT DISTINCT z.name, u.store
FROM Customer z, Purchase u, StorePrice v
WHERE z.cid = u.customer
AND u.store = v.store
AND v.price > 1000
```

Magda Balazinska - CSE 344, Fall 2012

3

Types of Views

- Virtual views**
 - Used in databases
 - Computed only on-demand – slow at runtime
 - Always up to date
- Materialized views**
 - Used in data warehouses
 - Pre-computed offline – fast at runtime
 - May have stale data
 - Indexes are materialized views

Magda Balazinska - CSE 344, Fall 2012

4

Customer(cid, name, city) StorePrice(store, price)
Purchase(customer, product, store)
Product(pname, price)

Query Modification

For each customer, find all the high end stores that they visit.

View:

```
CREATE VIEW StorePrice AS
SELECT x.store, y.price
FROM Purchase x, Product y
WHERE x.product = y.pname
```

Query:

```
SELECT DISTINCT z.name, u.store
FROM Customer z, Purchase u, StorePrice v
WHERE z.cid = u.customer
AND u.store = v.store
AND v.price > 1000
```

Magda Balazinska - CSE 344, Fall 2012

5

Customer(cid, name, city) StorePrice(store, price)
Purchase(customer, product, store)
Product(pname, price)

Query Modification

For each customer, find all the high end stores that they visit.

Modified query:

```
SELECT DISTINCT z.name, u.store
FROM Customer z, Purchase u,
(SELECT x.store, y.price
FROM Purchase x, Product y
WHERE x.pname = y.product) v
WHERE z.cid = u.customer
AND u.store = v.store
AND v.price > 1000
```

Magda Balazinska - CSE 344, Fall 2012

6

Customer(cid, name, city) StorePrice(store, price)
Purchase(customer, product, store)
Product(pname, price)

Query Modification

For each customer, find all the high end stores that they visit.

Modified and unnested query:

```
SELECT DISTINCT z.name, u.store
FROM Customer z, Purchase u,
     Purchase x, Product y
WHERE z.cid = u.customer
AND u.store = x.store
AND y.price > 1000
AND x.product = y.pname
```

Note that Purchase occurs twice. It has to be that way (why?).

Magda Balazinska - CSE 344, Fall 2012 7

Customer(cid, name, city) AcmePurchase(cid, cname, ..., price)
Purchase(customer, product, store)
Product(pname, price)

Further Virtual Views Optimizations

```
CREATE VIEW AcmePurchase AS
SELECT x.cid, x.name as cname, x.city, z.pid, z.name as pname, z.price
FROM Customer x, Purchase y, Product z
WHERE x.cid = y.customer and y.store = 'ACME' and y.product = z.pname
```

Query: `SELECT max(u.price) FROM AcmePurchase u` Find the highest prices at all ACME stores

View

```
SELECT max(z.price)
FROM Customer x, Purchase y, Product z
WHERE x.cid = y.customer and y.store = 'ACME' and y.product = z.pname
```

First rewrite. Can we further optimize?

Magda Balazinska - CSE 344, Fall 2012

Customer(cid, name, city) AcmePurchase(cid, name, ..., price)
Purchase(customer, product, store)
Product(pname, price)

Further Virtual Views Optimizations

```
CREATE VIEW AcmePurchase AS
SELECT x.cid, x.name as cname, x.city, z.pid, z.name as pname, z.price
FROM Customer x, Purchase y, Product z
WHERE x.cid = y.cid and y.store = 'ACME' and y.pid = z.pid
```

Query: `SELECT max(u.price) FROM AcmePurchase u` Find the highest prices at all ACME stores

View

```
SELECT max(z.price)
FROM Customer x, Purchase y, Product z
WHERE x.cid = y.customer and y.store = 'ACME' and y.product = z.pname
```

Correct if Purchase.customer is Not NULL and Foreign Key

9

Applications of Virtual Views

- Increased physical data independence. E.g.
 - Vertical data partitioning
 - Horizontal data partitioning
- Logical data independence. E.g.
 - Change schemas of base relations (i.e., stored tables)
- Security
 - View reveals only what the users are allowed to know

Magda Balazinska - CSE 344, Fall 2012 10

Vertical Partitioning

Resumes

SSN	Name	Address	Resume	Picture
234234	Mary	Huston	Clob1...	Blob1...
345345	Sue	Seattle	Clob2...	Blob2...
345343	Joan	Seattle	Clob3...	Blob3...
234234	Ann	Portland	Clob4...	Blob4...

T1

SSN	Name	Address
234234	Mary	Huston
345345	Sue	Seattle
...		

T2

SSN	Resume
234234	Clob1...
345345	Clob2...

T3

SSN	Picture
234234	Blob1...
345345	Blob2...

T2.SSN is a key and a foreign key to T1.SSN. Same for T3.SSN

11

T1(ssn,name,address) Resumes(ssn,name,address,resume,picture)
T2(ssn,resume)
T3(ssn,picture)

Vertical Partitioning

```
CREATE VIEW Resumes AS
SELECT T1.ssn, T1.name, T1.address,
       T2.resume, T3.picture
FROM T1,T2,T3
WHERE T1.ssn=T2.ssn and T1.ssn=T3.ssn
```

Magda Balazinska - CSE 344, Fall 2012 12

T1(ssn,name,address) Resumes(ssn,name,address,resume,picture)
 T2(ssn,resume)
 T3(ssn,picture)

Vertical Partitioning

```

SELECT address
FROM Resumes
WHERE name = 'Sue'
  
```

SELECT T1.address
 FROM T1, T2, T3
 WHERE T1.name = 'Sue'
 and T1.SSN=T2.SSN and T1.SSN = T3.SSN

Which of the tables T1, T2, T3 will be queried by the system ?

When do we use vertical partitioning ?

13

Vertical Partitioning Applications

- Advantages**
 - Speeds up queries that touch only a small fraction of columns
 - Single column can be compressed effectively, reducing disk I/O
- Disadvantages**
 - Updates are expensive!
 - Need many joins to access many columns
 - Repeated key columns add overhead

Hot trend today for data analytics: e.g., Vertica startup acquired by HP
They use a highly-tuned column-oriented data store AND engine

14

Horizontal Partitioning

Customers

SSN	Name	City
234234	Mary	Houston
345345	Sue	Seattle
345343	Joan	Seattle
234234	Ann	Portland
--	Frank	Calgary
--	Jean	Montreal

CustomersInHouston
 SSN Name City
 234234 Mary Houston

CustomersInSeattle
 SSN Name City
 345345 Sue Seattle
 345343 Joan Seattle

.....

Magda Balazinska - CSE 344, Fall 2012

15

CustomersInHouston(ssn,name,city)
 CustomersInSeattle(ssn,name,city)

Horizontal Partitioning

```

CREATE VIEW Customers AS
CustomersInHouston
UNION ALL
CustomersInSeattle
UNION ALL
...
  
```

Magda Balazinska - CSE 344, Fall 2012

16

CustomersInHouston(ssn,name,city)
 CustomersInSeattle(ssn,name,city)

Horizontal Partitioning

```

SELECT name
FROM Customers
WHERE city = 'Seattle'
  
```

Which tables are inspected by the system ?

Magda Balazinska - CSE 344, Fall 2012

17

CustomersInHouston(ssn,name,city)
 CustomersInSeattle(ssn,name,city)

Horizontal Partitioning

```

SELECT name
FROM Customers
WHERE city = 'Seattle'
  
```

Which tables are inspected by the system ?

All tables!
The systems doesn't know that CustomersInSeattle.city = 'Seattle'

Magda Balazinska - CSE 344, Fall 2012

18

CustomersInHouston(ssn,name,city)
CustomersInSeattle(ssn,name,city)

Customers(ssn,name,city)

Horizontal Partitioning

Better: remove CustomerInHouston.city etc

```

CREATE VIEW Customers AS
  (SELECT ssn, name, 'Houston' as city
   FROM CustomersInHouston)
 UNION ALL
  (SELECT ssn, name, 'Seattle' as city
   FROM CustomersInSeattle)
 UNION ALL
  ...

```

Magda Balazinska - CSE 344, Fall 2012
19

CustomersInHouston(ssn,name)
CustomersInSeattle(ssn,name)

Customers(ssn,name,city)

Horizontal Partitioning

```

SELECT name
FROM Customers
WHERE city = 'Seattle'

```

```

SELECT name
FROM CustomersInSeattle

```

Magda Balazinska - CSE 344, Fall 2012
20

Horizontal Partitioning Applications

- Performance optimization**
 - Especially for data warehousing
 - E.g. one partition per month
 - E.g. archived applications and active applications
- Distributed and parallel databases**
- Data integration**

Magda Balazinska - CSE 344, Fall 2012
21

Levels of Abstraction

Magda Balazinska - CSE 344, Fall 2012
22