

# Pig: High-level Procedural Language over MapReduce

# Matrix Addition

```
SELECT A.row, A.column, A.value + B.value
FROM A, B
WHERE A.row = B.row
      AND A.column = B.column
```

```
SELECT A.row, A.column, A.value + B.value
FROM A INNER JOIN B ON (
    A.row = B.row
    AND A.column = B.column
)
```

```
CREATE VIEW totalNumOfDocuments AS
SELECT count(frequency.docid) as count
FROM frequency
```

```
SELECT term_id f.term, docid f.docid, f.count *
log(N.count / (
    SELECT count(docid)
    FROM frequency f1
    WHERE f.term = f1.term)
)
FROM frequency f, totalNumOfDocument as N;
```

```
CREATE VIEW termdocs as
SELECT term, count(docid) as termdocs
FROM frequency
GROUP BY term;
```

```
CREATE VIEW totaldocs as
SELECT count(distinct docid) as countalldocs
FROM frequency;
```

```
SELECT a.term, a.docid, (
    a.count * log(b.termdocs /
    c.countalldocs)
) as tfidf
FROM frequency a, termdocs b, totaldocs c
WHERE a.docid = b.docid
AND a.term = b.term;
```

```
set @TotalNumberOfDocuments = (select COUNT(distinct doc_id) from frequency)

declare @DocumentCounts table
( term_id varchar(max),
  document_count int)

-- get the number of documents containing each term
insert into @DocumentCounts (term_id,document_count)
select term_id,COUNT(distinct doc_id)
from DATASCI250.dbo.frequency
group by term_id

-- Then calculate the TF-IDF, note that the below uses the Natural LOG
select doc_id,f.term_id,frequency * LOG(CAST(@TotalNumberOfDocuments as
Decimal(18,4))/CAST(Document_count as Decimal(18,4))) as TF_IDF
from frequency f
inner join @DocumentCounts d on f.term_id = d.term_id
```

```
select term_id, doc_id, frequency,
/* TF */
    (cast(frequency as float) /
    (select top 1 cast(frequency as float) as docFreq
     from [billhowe].[reuters_terms.csv] A
     where A.doc_id = Z.doc_id
     order by frequency desc
    ) *
/* IDF */
    (log(
        (select cast(count (distinct doc_id) as float)
         from [billhowe].[reuters_terms.csv]
        ) /
        (1 + (select cast(count (distinct doc_id) as float)
                 from [billhowe].[reuters_terms.csv] C
                 where Z.term_id = C.term_id )
        )
    )) as tfidf
from [billhowe].[reuters_terms.csv] Z
--order by doc_id, tfidf
```

# Reflection

- Why did we do this exercise?

# From the first lecture

What is Data Science about?

- 1) Preparing data for analysis  
(wrangling, cleaning, munging, transforming, integrating, ...)
- 2) Running some analysis (often a statistical model)
- 3) Interpreting the results and making decisions



## Key challenges for Data Preparation

- Data can be very large (Volume)
- Data can be very heterogeneous and weakly structured (Variety)
- Data may be coming in faster than you can handle it (Velocity)

## For Volume

- Key idea: use abstractions that allow scalable processing
  - Relational algebra and SQL
    - (not necessarily just databases)
  - MapReduce
  - Today: Pig

## For Variety

- Key idea: we want to know how to work with a variety of data types
  - Matrices
  - Relations
  - Graphs (some today)
  - Text
  - Images? (maybe)

## For Velocity

- We won't discuss this much in this class

# Key Challenges for Running Analyses


- Selecting the model
  - We will learn a toolbox of core techniques.
    - clustering (k-means)
    - optimization (forms of regression)
    - dimension reduction (multidimensional scaling)
- Running the model efficiently
  - *We will already know how to implement these models in existing systems*
    - *RA/SQL, MapReduce*
  - *We may have time to cover some new approaches (GraphLab)*

## Key Challenges for Interpreting Results

- Interpreting the results (convincing yourself)
- Communicating the results (convincing others)
- *We will focus on visualization*

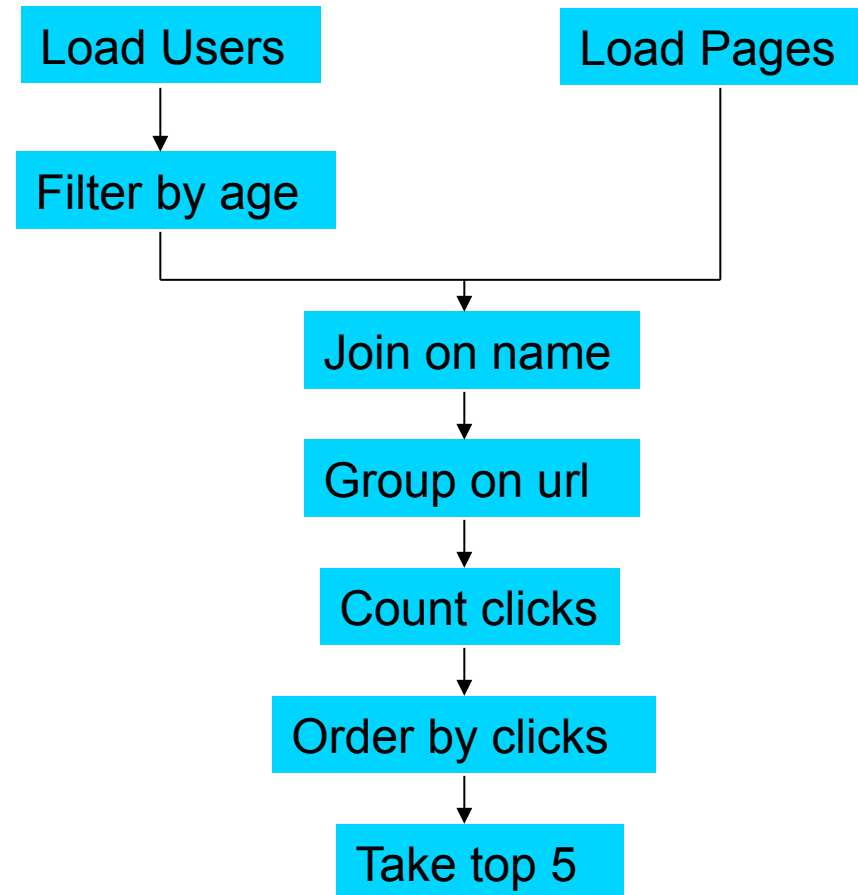
# What is Pig?

- An engine for executing programs on top of Hadoop
- It provides a language, Pig Latin, to specify these programs
- An Apache open source project <http://hadoop.apache.org/pig/>



## Why use Pig?

Suppose you have user data in one file, website data in another, and you need to find the top 5 most visited sites by users aged 18 - 25.





# In MapReduce

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.RecordReader;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.JobControl;
import org.apache.hadoop.mapred.lib.IdentityMapper;

public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String key = line.substring(0, firstComma);
            String value = line.substring(firstComma + 1);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("1" + value);
            oc.collect(outKey, outVal);
        }

        public static class LoadAndFilterUsers extends MapReduceBase
            implements Mapper<LongWritable, Text, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(firstComma + 1);
            int age = Integer.parseInt(value);
            if (age < 18 || age > 25) return;
            String key = line.substring(0, firstComma);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("2" + value);
            oc.collect(outKey, outVal);
        }

        public static class Join extends MapReduceBase
            implements Reducer<Text, Text, Text, Text> {

        public void reduce(Text key,
            Iterator<Text> iter,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // For each value, figure out which file it's from and
            // accordingly.
            List<String> first = new ArrayList<String>();
            List<String> second = new ArrayList<String>();

            while (iter.hasNext()) {
                Text t = iter.next();
                String value = t.toString();
                if (value.charAt(0) == '1')
                    first.add(value.substring(1));
                else second.add(value.substring(1));
            }

            reporter.setStatus("OK");
        }

        // Do the cross product and collect the values
        for (String s1 : first) {
            for (String s2 : second) {
                String outval = key + "," + s1 + "," + s2;
                oc.collect(null, new Text(outval));
                reporter.setStatus("OK");
            }
        }
    }

    public static class LoadJoined extends MapReduceBase
        implements Mapper<Text, Text, Text, LongWritable> {

        public void map(
            Text k,
            Text val,
            OutputCollector<Text, LongWritable> oc,
            Reporter reporter) throws IOException {
            // Find the url
            String line = val.toString();
            int firstComma = line.indexOf(',');
            int secondComma = line.indexOf(',', firstComma);
            String key = line.substring(firstComma, secondComma);
            // drop the rest of the record, I don't need it anymore,
            // just pass a 1 for the combiner/reducer to sum instead.
            Text outKey = new Text(key);
            oc.collect(outKey, new LongWritable(1L));
        }

        public static class ReduceURLs extends MapReduceBase
            implements Reducer<Text, LongWritable, WritableComparable,
                Writable> {

        public void reduce(
            Text key,
            Iterator<LongWritable> iter,
            OutputCollector<WritableComparable, Writable> oc,
            Reporter reporter) throws IOException {
            // Add up all the values we see
            long sum = 0;
            while (iter.hasNext()) {
                sum += iter.next().get();
                reporter.setStatus("OK");
            }
            oc.collect(key, new LongWritable(sum));
        }

        public static class LoadClicks extends MapReduceBase
            implements Mapper<WritableComparable, Writable, LongWritable,
                Text> {

        public void map(
            WritableComparable key,
            Writable val,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            oc.collect(((LongWritable)val), (Text)key);
        }

        public static class LimitClicks extends MapReduceBase
            implements Reducer<LongWritable, Text, LongWritable, Text> {

        int count = 0;
        public void reduce(
            LongWritable key,
            Iterator<Text> iter,
            OutputCollector<LongWritable, Text> oc,
            Reporter reporter) throws IOException {
            // Only output the first 100 records
            while (count < 100 && iter.hasNext()) {
                oc.collect(key, iter.next());
                count++;
            }
        }

        public static void main(String[] args) throws IOException {
            JobConf ip = new JobConf(MRExample.class);
            ip.setJobName("Load Pages");
            ip.setInputFormat(TextInputFormat.class);

            ip.setOutputKeyClass(Text.class);
            ip.setOutputValueClass(Text.class);
            ip.setMapperClass(LoadPages.class);
            FileInputFormat.addInputPath(ip, new
                Path("/user/gates/pages"));
            FileOutputFormat.setOutputPath(ip,
                new Path("/user/gates/tmp/indexed_pages"));
            ip.setNumReduceTasks(0);
            Job loadPages = new Job(ip);

            JobConf ifu = new JobConf(MRExample.class);
            ifu.setJobName("Load and Filter Users");
            ifu.setInputFormat(TextInputFormat.class);
            ifu.setOutputKeyClass(Text.class);
            ifu.setOutputValueClass(Text.class);
            ifu.setMapperClass(LoadAndFilterUsers.class);
            FileInputFormat.addInputPath(ifu, new
                Path("/user/gates/users"));
            FileOutputFormat.setOutputPath(ifu,
                new Path("/user/gates/tmp/filtered_users"));
            ifu.setNumReduceTasks(0);
            Job loadUsers = new Job(ifu);

            JobConf join = new JobConf(MRExample.class);
            join.setJobName("Join Users and Pages");
            join.setInputFormat(KeyValueTextInputFormat.class);
            join.setOutputKeyClass(Text.class);
            join.setOutputValueClass(Text.class);
            join.setMapperClass(IdentityMapper.class);
            join.setReducerClass(Join.class);
            FileInputFormat.addInputPath(join, new
                Path("/user/gates/tmp/indexed_pages"));
            FileInputFormat.addInputPath(join, new
                Path("/user/gates/tmp/filtered_users"));
            FileOutputFormat.setOutputPath(join, new
                Path("/user/gates/tmp/joined"));
            join.setNumReduceTasks(50);
            Job joinJob = new Job(join);
            joinJob.addDependingJob(loadPages);
            joinJob.addDependingJob(loadUsers);

            JobConf group = new JobConf(MRExample.class);
            group.setJobName("Group URLs");
            group.setInputFormat(KeyValueTextInputFormat.class);
            group.setOutputKeyClass(Text.class);
            group.setOutputValueClass(LongWritable.class);
            group.setOutputFormat(SequenceFileOutputFormat.class);
            group.setMapperClass(LoadJoined.class);
            group.setCombinerClass(ReduceURLs.class);
            group.setReducerClass(ReduceURLs.class);
            FileInputFormat.addInputPath(group, new
                Path("/user/gates/tmp/joined"));
            FileOutputFormat.setOutputPath(group, new
                Path("/user/gates/tmp/grouped"));
            group.setNumReduceTasks(50);
            Job groupJob = new Job(group);
            groupJob.addDependingJob(joinJob);

            JobConf top100 = new JobConf(MRExample.class);
            top100.setJobName("Top 100 sites");
            top100.setInputFormat(SequenceFileInputFormat.class);
            top100.setOutputKeyClass(LongWritable.class);
            top100.setOutputValueClass(Text.class);
            top100.setMapperClass(LoadClicks.class);
            top100.setMapperClass(LoadClicks.class);
            top100.setReducerClass(LimitClicks.class);
            top100.setCombinerClass(LimitClicks.class);
            FileInputFormat.addInputPath(top100, new
                Path("/user/gates/tmp/grouped"));
            FileOutputFormat.setOutputPath(top100, new
                Path("/user/gates/top100sitesforusers18to25"));
            top100.setNumReduceTasks(1);
            Job limit = new Job(top100);
            limit.addDependingJob(groupJob);

            JobControl jc = new JobControl("Find top 100 sites for users
                18 to 25");
            jc.addJob(loadPages);
            jc.addJob(loadUsers);
            jc.addJob(joinJob);
            jc.addJob(groupJob);
            jc.addJob(limit);
            jc.run();
        }
    }
}
```

170 lines of code, 4 hours to write

## In Pig Latin


```
Users = load 'users' as (name, age);  
Fltrd = filter Users by  
    age >= 18 and age <= 25;  
Pages = load 'pages' as (user, url);  
Jnd = join Fltrd by name, Pages by user;  
Grpd = group Jnd by url;  
Smmd = foreach Grpd generate group,  
    COUNT(Jnd) as clicks;  
Srted = order Smmd by clicks desc;  
Top5 = limit Srted 5;  
store Top5 into 'top5sites';
```

9 lines of code, 15 minutes to write

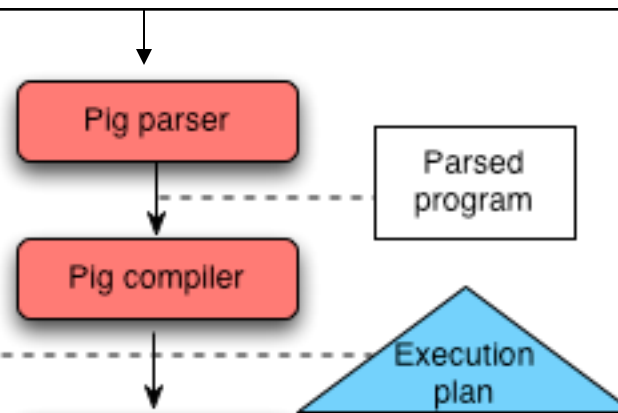
# Pig System Overview




Pig Latin  
program



```
A = LOAD 'file1' AS (sid,pid,mass,px:double);  
B = LOAD 'file2' AS (sid,pid,mass,px:double);  
C = FILTER A BY px < 1.0;  
D = JOIN C BY sid,  
        B BY sid;  
STORE g INTO 'output.txt';
```



## Pig Performance vs Map-Reduce



# Data Model

- “ Atom - simple atomic value (ie: number or string)
- “ Tuple
- “ Bag
- “ Map

# Data Model

---

- “ Atom
- “ Tuple - sequence of fields; each field any type
- “ Bag
- “ Map

$$\left( \text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

---

# Data Model

---

- “ Atom
- “ Tuple
- “ Bag - collection of tuples
  - “ Duplicates possible
  - “ Tuples in a bag can have different field lengths and field types
- “ Map

$$\left( \text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$

# Data Model

---

- “ Atom
- “ Tuple
- “ Bag
- “ Map - collection of key-value pairs
  - “ Key is an atom; value can be any type

$$\left( \text{'alice'}, \left\{ \begin{array}{l} (\text{'lakers'}, 1) \\ (\text{'iPod'}, 2) \end{array} \right\}, [\text{'age'} \rightarrow 20] \right)$$



# Speaking Pig Latin

---

## “ LOAD

- “ Input is assumed to be a bag (sequence of tuples)
- “ Can specify a deserializer with “USING”
- “ Can provide a schema with “AS”

```
newBag = LOAD 'filename'  
        <USING functionName() >  
        <AS (fieldName1, fieldName2,...)>;
```

```
Queries = LOAD 'query_log.txt'  
        USING myLoad()  
        AS (userID, queryString, timeStamp)
```



# Speaking Pig Latin

---

## “ FOREACH

- “ Apply some processing to each tuple in a bag
- “ Each field can be:
  - “ A fieldname of the bag
  - “ A constant
  - “ A simple expression (ie: f1+f2)
  - “ A predefined function (ie: SUM, AVG, COUNT, FLATTEN)
  - “ A UDF (ie: sumTaxes(gst, pst) )

```
newBag =  
    FOREACH bagName  
    GENERATE field1, field2, ...;
```



# Speaking Pig Latin

---

## “ FILTER

“ Select a subset of the tuples in a bag

```
newBag = FILTER bagName  
        BY      expression  ;
```

“ Expression uses simple comparison operators (==, !=, <, >, ...) and Logical connectors (AND, NOT, OR)

```
some_apples =  
    FILTER apples BY colour != 'red' ;
```

“ Can use UDFs

```
some_apples =  
    FILTER apples BY NOT isRed(colour) ;
```



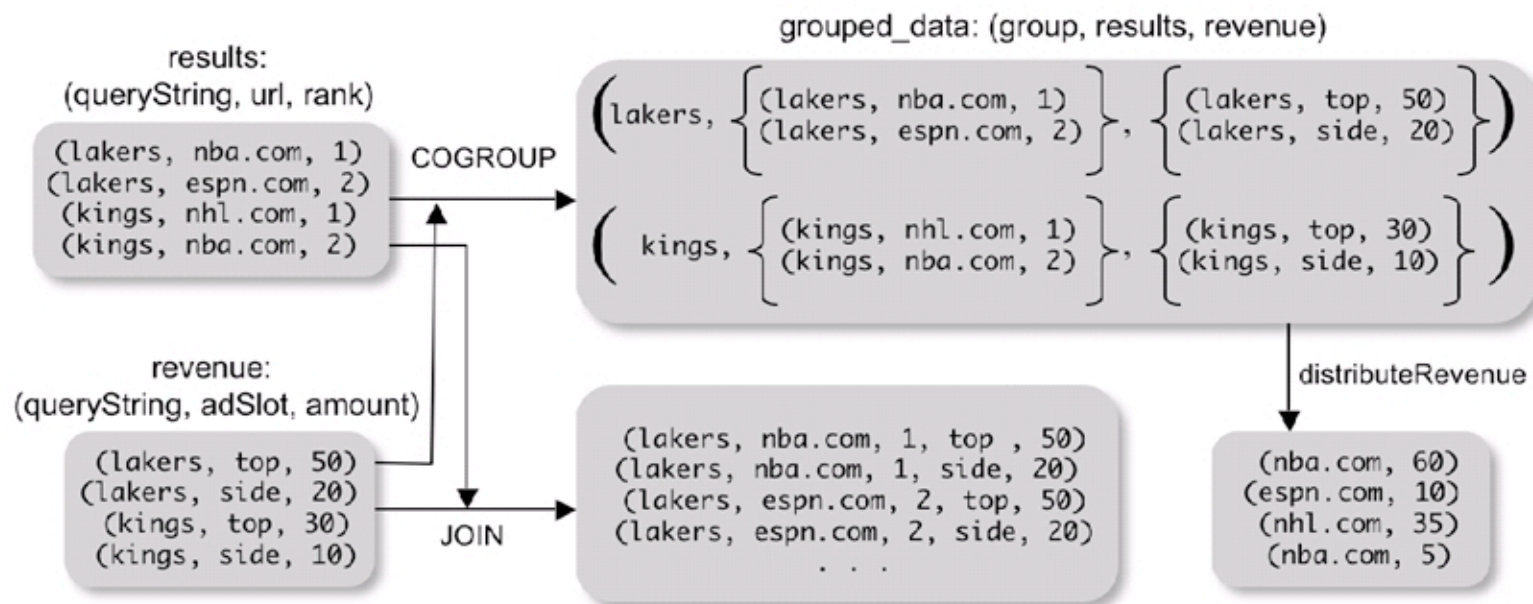
# Speaking Pig Latin

## “ COGROUP

“ Group two datasets together by a common attribute

“ Groups data into nested bags

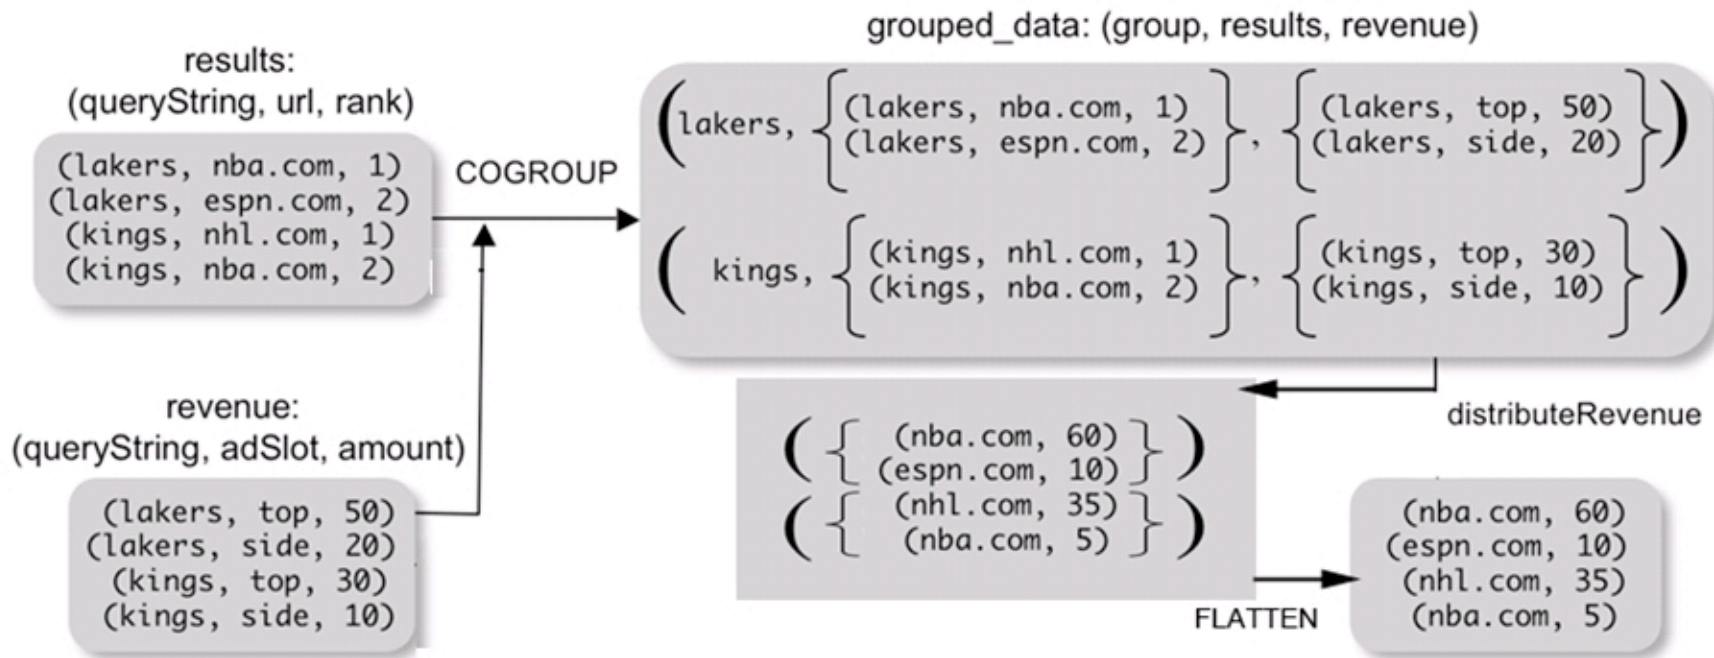
```
grouped_data = COGROUP results BY queryString,  
                           revenue BY queryString;
```



# Speaking Pig Latin

## “Why COGROUP and not JOIN?”

```
url_revenues =  
    FOREACH grouped_data GENERATE  
    FLATTEN(distributeRev(results, revenue));
```



# Speaking Pig Latin

---

## “Why COGROUP and not JOIN?

- “ May want to process nested bags of tuples before taking the cross product.
- “ Keeps to the goal of a single high-level data transformation per pig-latin statement.
- “ However, JOIN keyword is still available:

```
JOIN results BY queryString,  
    revenue BY queryString;
```



Equivalent

```
temp = COGROUP results BY queryString,  
    revenue BY queryString;  
join_result = FOREACH temp GENERATE  
    FLATTEN(results), FLATTEN(revenue);
```

# Speaking Pig Latin

---

## “ STORE (& DUMP)

“ Output data to a file (or screen)

```
STORE bagName INTO 'filename'  
<USING deserializer ()>;
```

## “ Other Commands (incomplete)

“ UNION - return the union of two or more bags

“ CROSS - take the cross product of two or more bags

“ ORDER - order tuples by a specified field(s)

“ DISTINCT - eliminate duplicate tuples in a bag

“ LIMIT - Limit results to a subset

# Compilation

---

- “ Pig system does two tasks:
  - “ Builds a Logical Plan from a Pig Latin script
    - “ Supports execution platform independence
    - “ No processing of data performed at this stage
  - “ Compiles the Logical Plan to a Physical Plan and Executes
    - “ Convert the Logical Plan into a series of Map-Reduce statements to be executed (in this case) by Hadoop Map-Reduce



# Compilation

---

## “ Building a Logical Plan

- “ Verify input files and bags referred to are valid
- “ Create a logical plan for each bag(variable) defined

# Compilation

---

## “ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```



Load(user.dat)

# Compilation

---

## “ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```




Load(user.dat)

# Compilation

---

## “ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```




# Compilation

---

## “ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```




# Compilation

---

## “ Building a Logical Plan Example


```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```



# Compilation

## “ Building a Logical Plan Example

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```




# Compilation

## “ Building a Physical Plan

```
A = LOAD 'user.dat' AS (name, age, city);  
B = GROUP A BY city;  
C = FOREACH B GENERATE group AS city,  
    COUNT(A);  
D = FILTER C BY city IS 'kitchener'  
    OR city IS 'waterloo';  
STORE D INTO 'local_user_count.dat';
```

Only happens when output is  
specified by STORE or DUMP






# Compilation

---

## “ Building a Physical Plan

- “ Step 1: Create a map-reduce job for each COGROUP



# Compilation

## “ Building a Physical Plan

“ Step 1: Create a map-reduce job for each COGROUP

“ Step 2: Push other commands into the map and reduce functions where possible

“ May be the case certain commands require their own map-reduce job (ie: ORDER needs separate map-reduce jobs)

